

TRAINING SESSION NAME

Firewall

Table of content

- [A. Definition](#)
 - [1. Concept](#)
 - [2. Configuration](#)
 - [2.1 Ruleset](#)
 - [2.2 Table](#)
 - [2.3 Chain](#)
 - [2.4 Rule](#)
 - [2.5 Set](#)
 - [2.6 Map](#)
- [B. Examples](#)
 - [1. Moving from iptables to nftables](#)
- [C. References](#)

A. Definition

1. Concept

- Nftables is a netfilter project that aims to replace the existing {ip,ip6,arp,eb} tables framework.
- Main differences with iptables:
 - + New syntax: *iptables* command line tool uses a getopt_long()-based parser where keys are always preceded by double minus, eg. --key or one single minus, eg. -p tcp. In contrast, nftables uses a compact syntax inspired by *tcpdump*.
 - + Tables and chains are fully configurable: *iptables* has multiple pre-defined tables and base chains, all of which are registered even if you only need one of them. With nftables there are no pre-defined tables or chains.
 - + Multiple actions
 - + Combine rules

2. Configuration

2.1 Ruleset

Definition

-The ruleset keyword is used to identify the whole set of tables, chains, etc. currently in place in kernel.

Command

```
{list | flush} ruleset [family]
```

2.2 Table

Definition

- Tables are containers for chains, sets and stateful objects.
- They are identified by their address family and their name.
- Address families determine the type of packets which are processed.

Address families	Description
ip	IPv4 address family.
ip6	IPv6 address family.
inet	Internet (IPv4/IPv6) address family.
arp	ARP address family, handling IPv4 ARP packets.
bridge	Bridge address family, handling packets which traverse a bridge device.
netdev	Netdev address family, handling packets on ingress and egress.

Command

```
{add | create} table [family] table [ {comment comment ;} { flags 'flags ; }]  
{delete | destroy | list | flush} table [family] table  
list tables [family]  
delete table [family] handle handle  
destroy table [family] handle handle
```

2.3 Chain

Definition

- Chains are containers for rules. They exist in two kinds, base chains and regular chains.
- A base chain is an entry point for packets from the networking stack, a regular chain may be used as

jump target and is used for better rule organization.

- For base chains, **type**, **hook** and **priority** parameters are mandatory.

Type	Families	Hooks	Description
filter	all	all	Standard chain type to use in doubt.
nat	ip, ip6, inet	prerouting, input, output, postrouting	Chains of this type perform Native Address Translation based on conntrack entries. Only the first packet of a connection actually traverses this chain - its rules usually define details of the created conntrack entry (NAT statements for instance).
route	ip, ip6	output	If a packet has traversed a chain of this type and is about to be accepted, a new route lookup is performed if relevant parts of the IP header have changed. This allows one to e.g. implement policy routing selectors in nftables.

Hook	Families	Description
prerouting	ip, ipv6, inet, bridge	All packets entering the system are processed by the prerouting hook. It is invoked before the routing process and is used for early filtering or changing packet attributes that affect routing.
input	ip, ipv6, inet, bridge, arp	Packets delivered to the local system are processed by the input hook.
forward	ip, ipv6, inet, bridge	Packets forwarded to a different host are processed by the forward hook.
output	ip, ipv6, inet, bridge, arp	Packets sent by local processes are processed by the output hook.
postrouting	ip, ipv6, inet, bridge	All packets leaving the system are processed by the postrouting hook.
ingress	ip, ipv6, inet, bridge, netdev	All packets entering the system are processed by this hook. It is invoked before layer 3 protocol handlers, hence before the prerouting hook, and it can be used for filtering and policing. Ingress is only available for Inet family (since Linux kernel 5.10).



egress	netdev	All packets leaving the system are processed by this hook. It is invoked after layer 3 protocol handlers and before tc egress. It can be used for late filtering and policing.
--------	--------	---

- Standard priority names, family and hook compatibility matrix

Name	Value	Families	Hooks
raw	-300	ip, ip6, inet	all
mangle	-150	ip, ip6, inet	all
dstnat	-100	ip, ip6, inet	prerouting
filter	0	ip, ip6, inet, arp, netdev	all
security	50	ip, ip6, inet	all
srcnat	100	ip, ip6, inet	postrouting

- Standard priority names and hook compatibility for the bridge family

Name	Value	Hooks
dstnat	-300	prerouting
filter	-200	all
out	100	output
srcnat	300	postrouting

Command

```
{add | create} chain [family] table chain [{ type type hook hook [device device]  
priority priority ; [policy policy ;] [comment comment ;] }]
```

```
{delete | destroy | list | flush} chain [family] table chain  
list chains [family]  
delete chain [family] table handle handle  
destroy chain [family] table handle handle  
rename chain [family] table chain newname
```

2.4 Rule

Command

```
{add | insert} rule [family] table chain [handle handle | index index] statement  
... [comment comment]  
replace rule [family] table chain handle handle statement ... [comment comment]  
{delete | reset} rule [family] table chain handle handle  
destroy rule [family] table chain handle handle  
reset rules [family]  
reset rules table [family] table  
reset rules chain [family] table [chain]
```

2.5 Set

Definition

- Set consist of one or more elements, separated by commas, enclosed by curly brace.
- Anonymous sets are embedded in rules and cannot be updated, you must delete and re-add the rule.
- Named sets can be updated, and can be typed and flagged.

Keyword	Description	Type
type	data type of set elements	string: ipv4_addr, ipv6_addr, ether_addr, inet_proto, inet_service, mark
typeof	data type of set element	expression to derive the data type from
flags	set flags	string: constant, dynamic, interval, timeout
timeout	time an element stays in the set, mandatory if set is added to from the packet path (ruleset)	string, decimal followed by unit. Units are: d, h, m, s
gc-interval	garbage collection interval, only available when timeout or flag timeout are active	string, decimal followed by unit. Units are: d, h, m, s



elements	elements contained by the set	set data type
size	maximum number of elements in the set, mandatory if set is added to from the packet path (ruleset)	unsigned integer (64 bit)
policy	set policy	string: performance [default], memory
auto-merge	automatic merge of adjacent/overlapping set elements (only for interval sets)	

Command

```
add set [family] table set { type type | typeof expression ; [flags flags ;]
[timeout timeout ;] [gc-interval gc-interval ;] [elements = { element[, ...] } ;]
[size size ;] [comment comment ;] [policy 'policy ;] [auto-merge ;] }
{delete | destroy | list | flush} set [family] table set
list sets [family]
delete set [family] table handle handle
{add | delete | destroy } element [family] table set { element[, ...] }
```

2.6 Map

Definition

- Maps store data based on some specific key used as input. They are uniquely identified by a user-defined name and attached to tables.

Command

```
add map [family] table map { type type | typeof expression [flags flags ;]
[elements = { element[, ...] } ;] [size size ;] [comment comment ;] [policy
'policy ;] }
{delete | destroy | list | flush} map [family] table map
list maps [family]
```

B. Examples

Filter

Filter config

```
#!/usr/sbin/nft -f
flush ruleset

table ip myip {
    chain web {
        tcp dport http accept
```

```
        tcp dport 8080 accept
    }
    chain myinput {
        type filter hook input priority 0; policy drop;
        ip saddr 192.168.56.112 accept
        ip saddr 10.0.3.0/24 jump web
    }
}

table bridge mybridge {
    chain myoutput {
        type filter hook output priority 0; policy drop;
        ether saddr 08:00:27:ef:94:c3 accept
    }
}

table ip mynat {
    chain myprerouting {
        type nat hook prerouting priority dstnat;
        tcp dport { ssh, http } dnat to 10.0.3.1
    }
    chain mypostrouting {
        type nat hook postrouting priority srcnat;
        ip daddr 10.0.3.1 masquerade
    }
}
```

Rule

Rule command

```
$ nft add table ip myip
$ nft add chain ip myip myinput { type filter hook input priority 0; policy drop;
}
$ nft add rule ip myip myinput ip saddr 192.168.56.112 accept
$ nft add rule ip nat prerouting dnat tcp dport map { 80 : 192.168.1.100, 8888 :
192.168.1.101 }
```

1. Moving from iptables to nftables

- Translate by command

```
$ iptables-translate -A INPUT -p tcp --dport 22 -j ACCEPT
nft add rule ip filter INPUT tcp dport 22 counter accept
```

- Translate by whole ruleset

```
$ iptables-save > save.txt
$ cat save.txt
# Generated by iptables-save v1.6.0 on Sat Dec 24 14:26:40 2016
*filter
:INPUT ACCEPT [5166:1752111]
:FORWARD ACCEPT [0:0]
:OUTPUT ACCEPT [5058:628693]
```



```
-A FORWARD -p tcp -m tcp --dport 22 -j ACCEPT
COMMIT
# Completed on Sat Dec 24 14:26:40 2016
$ iptables-restore-translate -f save.txt
# Translated by iptables-restore-translate v1.6.0 on Sat Dec 24 14:26:59 2016
add table ip filter
add chain ip filter INPUT { type filter hook input priority 0; }
add chain ip filter FORWARD { type filter hook forward priority 0; }
add chain ip filter OUTPUT { type filter hook output priority 0; }
add rule ip filter FORWARD tcp dport 22 counter accept
```

C. References

No	Info	Link/file/name of ebook
1	nftables definition	https://wiki.archlinux.org/title/nftables

